

SERVICE MESHES

Craig Dillon

MSc DevOps

x00205790@mytudublin.ie



WHAT ARE SERVICE MESHES?

“Service meshes provide policy-based, network services for network-connected workloads by enforcing desired behavior of the network in the face of constantly changing conditions and topology.”

WHAT DOES THAT REALLY MEAN?

- Instead of adding logic to application code to handle communication, metrics and security functions between microservices, offload these workloads to a service mesh for management.
- This ensures a consistent internal network configuration across microservices, reducing development overhead and providing resiliency.






WHY SERVICE MESHES?

- First there were monoliths
- Eventually evolved into microservices where applications are now broken down into several microservice components
- Presented a new set of problems with service communication
- A service mesh is a means to allow reliable, secure and observable communication between these components

PRINCIPLES OF SERVICE MESHES



- Observability
- Routing
- Automatic scaling
- Separation of duties
- Trust
- Automatic service registration and discovery
- Resilience



COMPONENTS OF A SERVICE MESH

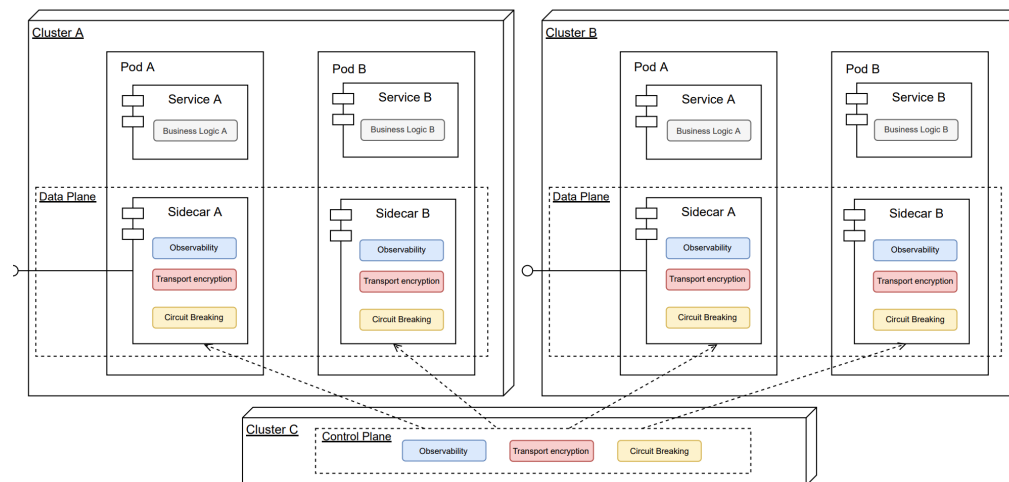
- Control plane: The Control Plane handles service discovery, metric collection, certificate management and configuration of the Data Plane.
- Data plane: Handles ingress and egress traffic for microservices, message processing, encryption enforcement and metric collection.

SUBCOMPONENTS OF A SERVICE MESH

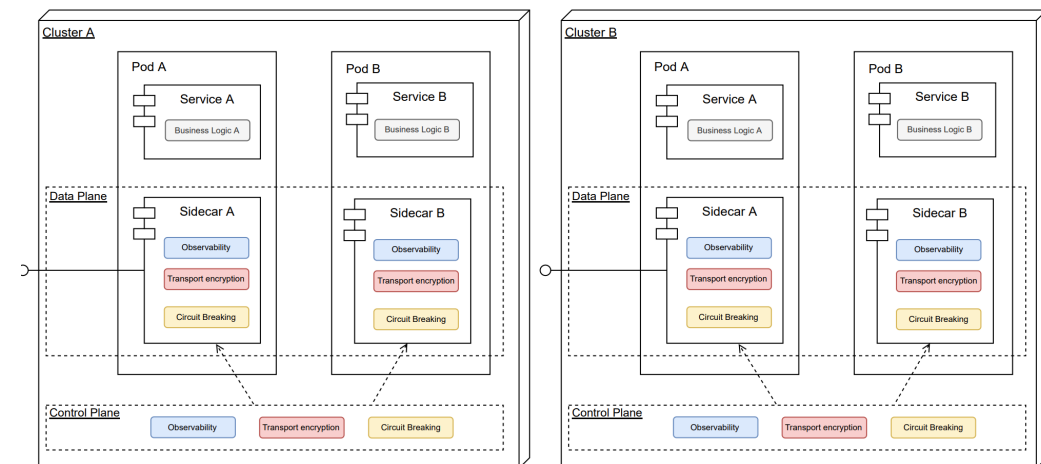
- Data Plane:
 - Proxy: This is usually in the form of a sidecar container running alongside the microservice in each pod.
- Control Plane:
 - Traffic manager – Performs service discovery and when a change is detected notifies proxies to modify routing as needed.
 - Security component – Encrypts traffic between microservices, often using *mutual Transport Layer Security* (mTLS) in a transparent manner without making changes to the application.
 - Configuration component – Responsible for validating and distributing configurations to components of the service mesh.



ARCHITECTURE OF SERVICE MESHES



Service Mesh with Sidecar



Multi Cluster Service Mesh

A network diagram featuring several colored nodes (red, green, blue) connected by black lines. The nodes are arranged in a non-linear fashion, with some nodes having multiple connections. The lines are thin and black, creating a web-like structure. The background is white.

DRAWBACKS AND CHALLENGES

- Increased complexity
- Performance overhead
- Operational overhead
- Vendor lock-in
- Learning curve



- Largest fashion E-commerce company in Japan
- Initially deployed as an on-premises monolithic platform in 2014
- In 2018 moved to a hybrid K8s platform with an API Gateway added in 2020 which added development overheads
- 2021 saw the introduction of Istio as a service mesh



- Xbox Cloud launched in 2019 and used an in-house developed mTLS (mutual TLS) to ensure secure communication between services
- As the service grew the development and maintenance overhead also grew
- Moved to Linkerd as the service mesh solution
- Kustomize was used to configure Linkerd and other components
- Service is built on AKS (Azure Kubernetes Service) and Linkerd works in conjunction with Flagger to allow canary deployments
- Massively reduced engineering time, improved reliability, saved thousands per month and reduced latency by 100ms

EMERGING TRENDS

- Moving service mesh functions to eBPF, Cilium is currently the leading service mesh applying this method
- Kuma is a solution that is designed to be environment agnostic and supports multi-zone, multi-cluster and multi-cloud deployments
- Traefik Mesh is designed to be a lightweight, non-invasive service mesh meaning it runs a dedicated pod on each node to handle routing. Services are added manually rather than via service discovery.

COMPARISON OF SERVICE MESHES



LINKERD



Istio



HashiCorp
Consul



cilium

Sidecar Proxy	Yes	Yes	Yes	Optional
Per node agent	No	No	Yes	Yes
Multicluster support	Yes	Yes	Yes	Yes
Deployment method	Helm	Helm	Helm	Helm
Ingress controller	Any	Istio ingress or Istio gateway	Envoy & ambassador	Cilium Ingress, compatible with others
Other considerations			Requires Consul server for quorum operations	Is a Kubernetes network component rather than service



CRITICAL SUMMARY

- Streamlined communication between microservices
- Does not replace operational overhead but potentially reduces it
- Not all service meshes are created equal
- Rapidly evolving ecosystem
- Overall provides a holistic means of managing networking between microservices

REFERENCES

- Calcote, L. & Butcher, Z. (2019), Istio: Up and running: Using a service mesh to connect, secure, control, and observe, O'Reilly Media.
- Duarte Maia, J. T. & Figueiredo Correia, F. (2022), Service mesh patterns, in 'Proceedings of the 27th European Conference on Pattern Languages of Programs', pp. 1–12.
- Kawasaki, Y. (2022), 'Accelerating zozotown modernization with istio'. URL: <https://istio.io/latest/about/case-studies/zozo/> (accessed 18/03/2024)
- Khatri, A. & Khatri, V. (2020), Mastering Service Mesh: Enhance, secure, and observe cloud-native applications with Istio, Linkerd, and Consul, Packt Publishing Ltd.
- Koschel, A., Bertram, M., Bischof, R., Schulze, K., Schaaf, M. & Astrova, I. (2021), A look at service meshes, in '2021 12th International Conference on Information, Intelligence, Systems & Applications (IISA)', IEEE, pp. 1–8.
- Li, W., Lemieux, Y., Gao, J., Zhao, Z. & Han, Y. (2019), Service mesh: Challenges, state of the art, and future research opportunities, in '2019 IEEE International Conference on Service-Oriented System Engineering (SOSE)', IEEE, pp. 122–1225.
- Wodajie, A. & Voss, C. (2022), 'Service mesh at scale: How xbox cloud gaming secures 22k pods with linkerd'. URL: <https://www.cncf.io/blog/2022/05/10/service-mesh-at-scale-how-xbox-cloud-gaming-secures-22k-pods-with-linkerd> (accessed 23/03/2024)
- Zhu, X., She, G., Xue, B., Zhang, Y., Zhang, Y., Zou, X. K., Duan, X., He, P., Krishnamurthy, A., Lentz, M. et al. (2022), 'Dissecting service mesh overheads', arXivpreprint arXiv:2207.00592 .3